

Comparing Deep Learning Transformers To Naive Bayes and Random Forest Models in the Context of Emotion Detection

Richard Hoehn

Middle Tennessee State University

```
In [1]: # Imports of main Libs
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# SKlearn
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix

# BERT - Huggingface
from transformers import BertModel
from transformers import BertTokenizer
from transformers import BertForSequenceClassification

# PyTorch Libs
import torch
from torch.utils.data import TensorDataset, DataLoader
import torch.nn.functional as F
from torchmetrics import Accuracy
from torch.optim.lr_scheduler import StepLR
import torch.optim as optim

# Lightning Libs - Primarily for the Trainer()
import pytorch_lightning as pl
from pytorch_lightning import Trainer
from pytorch_lightning.loggers import CSVLogger
from pytorch_lightning.callbacks.early_stopping import EarlyStopping

# Plotting and Displaying Setup
%matplotlib inline
pd.set_option('display.max_columns', 500) # Allows for up to 500 columns to be displayed when viewing a dataframe
pd.set_option('display.max_info_columns', 250) # Allows for up to 250 columns to be displayed when viewing a dataframe
plt.style.use('seaborn-v0_8') # A style that can be used for plots - see style reference above

# Set Device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device
```

```
Out[1]: device(type='cuda')
```

Data Loading

```
In [2]: # Load csv into dataframe
df = pd.read_csv('pd_en_translated.csv', index_col = None, header=0)

# Display Some details of the Dataset
print("\n*** Emotion Tagged Sentences ***")
display(df.head(5))

*** Emotion Tagged Sentences ***
```

	emotion_en	sentence_en	emotion_de	sentence_de
0	sadness	sorry mum, will treat you well from now on	Traurigkeit	Entschuldigung, Mama, werde Sie von nun an gut...
1	hate	i hate snoring. remind me if my future husband...	Ekel	Ich hasse es zu schnarchen.Erinnern Sie mich d...
2	surprise	And I end up in privilege... Oh well, at least...	Überraschung	Und ich lande im Privileg ... na ja, zumindest...
3	worry	Tasha's really bad haircut. She's being treate...	Angst	Tasha ist wirklich schlechter Haarschnitt.Sie ...
4	sadness	I miss my mom.. "May angels lead you in&...	Traurigkeit	Ich vermisse meine Mutter.

```
In [3]: # Drop Columns Not Needed
df.drop(["emotion_de", "sentence_de"], axis=1, inplace=True)
```

Plot Pie Chart of Emotion Distribution

```
In [4]: # Group by Emotions
emotion_count = df["emotion_en"].value_counts()

# Make sure we use the entire width
plt.figure(figsize=(20, 10))

# Custom function to show both count and percentage on the pie chart
def autopct_format(values):
    def my_format(pct):
        total = sum(values)
        val = int(round(pct*total/100.0))
        return '{v:d} ({p:.0f}%)'.format(p=pct,v=val)
    return my_format

# Group by Emotions
emotion_count = df["emotion_en"].value_counts()

# First - Raw Pie Chart
plt.subplot(121)
plt.pie(emotion_count,
        labels = emotion_count.index,
        autopct = autopct_format(emotion_count)
        )

# Setup Title
plt.title(f'Distribution of Emotions by {len(df)} rows',
        fontweight='bold',
        fontsize='18',
        horizontalalignment='center')

# Remove rows where 'emotion_en' is 'anger' and 'enthusiasm'
df = df[df['emotion_en'] != 'anger']
df = df[df['emotion_en'] != 'enthusiasm']


# Group by Emotions
emotion_count = df["emotion_en"].value_counts()

# Second - Updated Pie Chart
plt.subplot(122)
plt.pie(emotion_count,
        labels = emotion_count.index,
        autopct = autopct_format(emotion_count)
        )


# Setup Title
plt.title(f'Distribution of Emotions by {len(df)} rows',
        fontweight='bold',
        fontsize='18',
        horizontalalignment='center')

# Save & Show the Plot
plt.savefig('fig_distribution.png')
plt.show()
```

Distribution of Emotions by 1500 rows



Distribution of Emotions by 1464 rows



Review Max Length and Distttribution of Sentence Counts

```
In [5]: df["word_count"] = df["sentence_en"].str.split().apply(len)

# Setting the Style of Seaborn
sns.set(style='dark')

# Make sure we use the entire width
plt.figure(figsize=(18, 8))

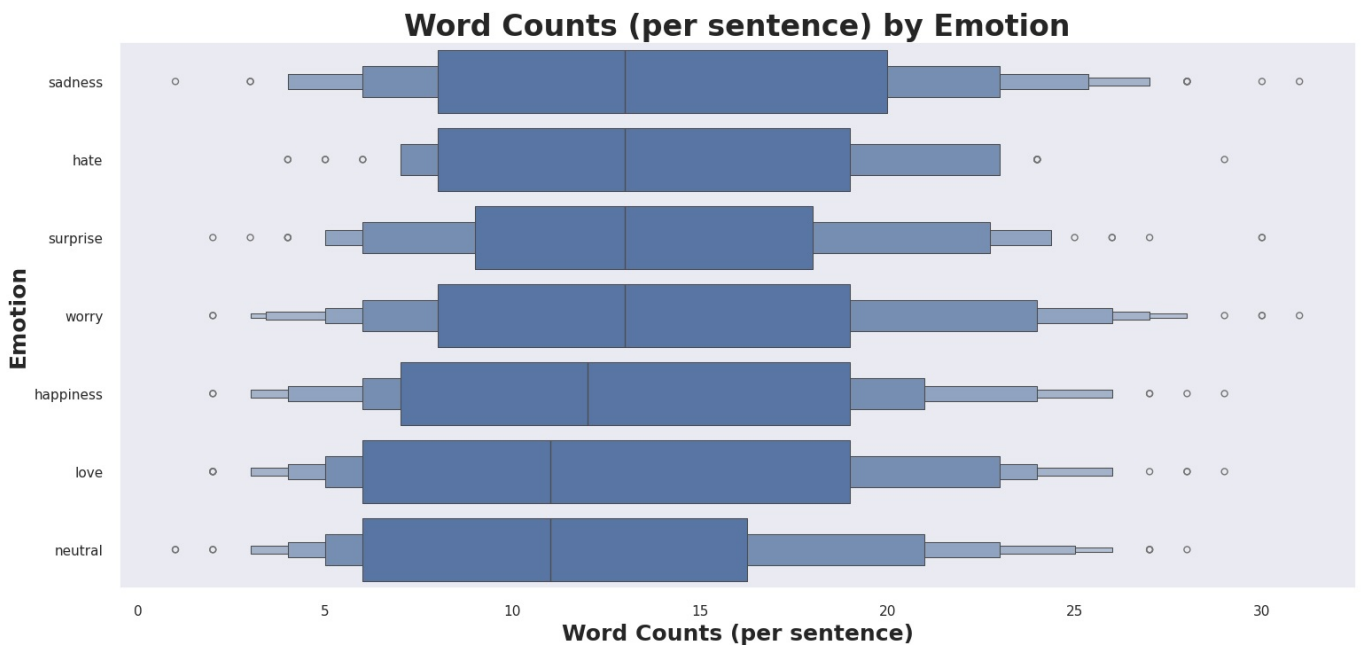
sns.boxenplot(
    data = df,
    x = "word_count",
    y = "emotion_en")

# Setup Title, and Axis Lables
plt.title('Word Counts (per sentence) by Emotion',
          fontweight='bold',
          fontsize='24',
          horizontalalignment='center')

plt.xlabel('Word Counts (per sentence)',
           fontweight='bold',
           fontsize='18',
           horizontalalignment='center')

plt.ylabel('Emotion',
          fontweight='bold',
          fontsize='18')

# Save & Show the Plot
plt.savefig('fig_word_counts.png')
plt.show()
```



Label Encoding of Emotions

```
In [6]: # Label Encoding the Emotions
label_encoder = LabelEncoder()

# Encode labels in column 'emotion_en'
df['label'] = label_encoder.fit_transform(df['emotion_en'])

# Show Updated Dataframe
display(df.head(10))
```

	emotion_en	sentence_en	word_count	label
0	sadness	sorry mum, will treat you well from now on	9	4
1	hate	i hate snoring. remind me if my future husband...	15	1
2	surprise	And I end up in privilege... Oh well, at least...	15	5
3	worry	Tasha's really bad haircut. She's being treate...	18	6
4	sadness	I miss my mom.. "May angels lead you in&...	9	4
5	surprise	@laurenceobrien Thank you xo	4	5
6	surprise	The chicken noodle soup I made for lunch to fe...	21	5
7	worry	Being dragged round Ikea this morning Bad times!	8	6
8	happiness	seniors done 5 more days!! woohoo!! going out...	11	0
9	worry	...wait, I lied.	3	6

```
In [7]: # Setup / Instantiate the Tokenizer
tokenizer = BertTokenizer.from_pretrained("bert-base-uncased", do_lower_case=True)

# Show Detail on the BERT Tokenizer
print(f'Vocab size: {tokenizer.vocab_size}')
print(f'Max length: {tokenizer.model_max_length}')
print(f'Tokenizer Model Input: {tokenizer.model_input_names}')
```

Vocab size: 30522

Max length: 512

Tokenizer Model Input: ['input_ids', 'token_type_ids', 'attention_mask']

Encoding / Tokenization Stage

```
In [8]: # Example of Tokenizer
print('Encoded text')
encoded_text = tokenizer("Hello World!",
                        add_special_tokens = True,
                        max_length = 10,
                        truncation=True,
                        padding=True)

print(encoded_text, '\n')
```

```
/opt/conda/lib/python3.11/site-packages/transformers/tokenization_utils_base.py:2618: FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future version, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use `padding='max_length'` to pad to a max length. In this case, you can give a specific length with `max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of the model (e.g. 512 for Bert).
  warnings.warn(
```

```
Tokens:  
-> {'input_ids': [101, 1030, 5532, 1035, 2732, 2683, 2629, 2821, 2061, 2017, 2113, 2129, 1045, 2514, 2059, 4365,  
4387, 2005, 2924, 1997, 2637, 2699, 2000, 2191, 2009, 2614, 2066, 1045, 2106, 2009, 1012, 2054, 1037, 1038, 1012  
, 1012, 1012, 102, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], 'token_type_id'  
s': [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0,  
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]}
```

Emotion:
-> worry

Unique Classes: 7

1171

1171

146

146

147

147

Build Datasets into Tensors

```
In [11]: # Create Tensors for Training and Testing

# Training Data
train_input_ids_tensors = torch.tensor([d['input_ids'] for d in X_train])
train_attention_mask_tensors = torch.tensor([d['attention_mask'] for d in X_train])
train_labels_tensors = torch.tensor(y_train)

# Validation Data
val_input_ids_tensors = torch.tensor([d['input_ids'] for d in X_val])
val_attention_mask_tensors = torch.tensor([d['attention_mask'] for d in X_val])
val_labels_tensors = torch.tensor(y_val)

# Test Data
test_input_ids_tensors = torch.tensor([d['input_ids'] for d in X_test])
test_attention_mask_tensors = torch.tensor([d['attention_mask'] for d in X_test])
test_labels_tensors = torch.tensor(y_test)
```

Creating the DataLoaders for the Trainer()

```
In [12]: # Select a batch size for training. For fine-tuning BERT on a specific task, the authors recommend a batch size
BATCH_SIZE = 10
NUM_WORKERS = 3

train_data = TensorDataset(train_input_ids_tensors,
                           train_attention_mask_tensors,
                           train_labels_tensors)

train_dataloader = DataLoader(train_data,
                              batch_size = BATCH_SIZE,
                              num_workers = NUM_WORKERS,
                              shuffle = True)

val_data = TensorDataset(val_input_ids_tensors,
                         val_attention_mask_tensors,
                         val_labels_tensors)

val_dataloader = DataLoader(val_data,
                            batch_size = BATCH_SIZE,
                            num_workers = NUM_WORKERS,
                            shuffle = False)

test_data = TensorDataset(test_input_ids_tensors,
                          test_attention_mask_tensors,
                          test_labels_tensors)

test_dataloader = DataLoader(test_data,
                             batch_size = BATCH_SIZE,
                             num_workers = NUM_WORKERS,
                             shuffle = False)
```

BERT Model Design

This model inherits from the `pl.LightningModule` to allow for the `Trainer()` to handle learning rate and early stopping.

```
In [13]: class BertClassifier(pl.LightningModule):
    def __init__(self, n_classes: int, learning_rate: float):
        super().__init__()
        self.bert = BertForSequenceClassification.from_pretrained("bert-base-uncased", num_labels=n_classes)
        self.learning_rate = learning_rate
        self.accuracy = Accuracy(num_classes = n_classes, task='multiclass')

    def forward(self, input_ids, attention_mask, labels=None):
        output = self.bert(input_ids, attention_mask=attention_mask, labels=labels)
        return output
```

```

def predict(self, input_ids, attention_mask):
    self.eval() # Set the model to evaluation mode
    with torch.no_grad(): # Disable gradient calculation
        outputs = self.bert(input_ids, attention_mask=attention_mask)
        logits = outputs.logits
        probs = F.softmax(logits, dim=1) # Convert logits to probabilities
    return probs

def training_step(self, batch, batch_idx):
    input_ids, attention_mask, labels = batch
    output = self.forward(input_ids, attention_mask, labels)

    # Get Loss
    loss = output.loss
    self.log('train_loss', loss, on_step=False, on_epoch=True)

    # Get Acc
    logits = output.logits
    preds = torch.argmax(logits, dim=1)
    acc = self.accuracy(preds, labels)
    self.log('train_acc', acc, on_step=False, on_epoch=True)

    # Get Learning Rate (for Logging)
    lr = self.trainer.optimizers[0].param_groups[0]['lr']
    self.log('lr', lr, on_step=False, on_epoch=True)

    return loss

def validation_step(self, batch, batch_idx):
    input_ids, attention_mask, labels = batch
    output = self.forward(input_ids, attention_mask, labels)

    # Get Loss
    loss = output.loss
    self.log('val_loss', loss, on_step=False, on_epoch=True)

    # Get Acc
    logits = output.logits
    preds = torch.argmax(logits, dim=1)
    acc = self.accuracy(preds, labels)
    self.log('val_acc', acc, on_step=False, on_epoch=True)

    return loss

def test_step(self, batch, batch_idx):
    input_ids, attention_mask, labels = batch
    output = self.forward(input_ids, attention_mask, labels)

    # Get Loss
    loss = output.loss
    self.log('test_loss', loss, on_step=False, on_epoch=True)

    # Get Acc
    logits = output.logits
    preds = torch.argmax(logits, dim=1)
    acc = self.accuracy(preds, labels)
    self.log('test_acc', acc, on_step=False, on_epoch=True)

    return loss

def configure_optimizers(self):
    # Define Optimizer
    optimizer = torch.optim.Adam(self.parameters(), lr=self.learning_rate)

    # Define LR scheduler
    scheduler = {
        'scheduler': StepLR(optimizer,
                             step_size = 1, # Drop Each Epoch
                             gamma = 0.95), # 5% Drop in LR
        'interval': 'epoch', # or 'step' for step-wise
        'frequency': 2,
    }
    # Return the Optimizer and Scheduler in Dictionary List
    return [optimizer], [scheduler]

```

```

In [14]: # Settings / Configurations
NUM_EPOCHS = 10
NUM_CLASSES = n_classes
L_RATE = 0.000005
LOG_DIR = "logs"
LOG_NAME = "Classifier"
ES_PATIENCE = 3

```

```

# Instantiate the Logger
csv_logger = CSVLogger(name = LOG_NAME, save_dir = LOG_DIR)

# Instantiate Model with Class Count
model = BertClassifier(n_classes = NUM_CLASSES, learning_rate = L_RATE)

# Initialize the EarlyStopping callback
early_stopping = EarlyStopping(
    monitor = 'val_loss',      # Monitor validation loss for early stopping
    patience = ES_PATIENCE,    # Number of epochs to wait for improvement before stopping
    verbose = True,            # Print messages when stopping
    mode = 'min'               # The direction is automatically inferred if not set but it's good practice to inc
)

# Setup Trainer with Logger Included
trainer = pl.Trainer(logger = csv_logger,
                    max_epochs = NUM_EPOCHS,
                    enable_progress_bar = True,
                    callbacks=[early_stopping])

# Start Training / Fine-Tuning the Model
trainer.fit(model, train_dataloader, val_dataloader)

```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']
 You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
 GPU available: True (cuda), used: True
 TPU available: False, using: 0 TPU cores
 IPU available: False, using: 0 IPUs
 HPU available: False, using: 0 HPUs
 LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

	Name	Type	Params
0	bert	BertForSequenceClassification	109 M
1	accuracy	MulticlassAccuracy	0

109 M Trainable params
 0 Non-trainable params
 109 M Total params
 437.950 Total estimated model params size (MB)
 SLURM auto-requeueing enabled. Setting signal handlers.

Sanity Checking: | | 0/? [00:00<?, ?it/s]

Training: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Metric val_loss improved. New best score: 1.779

Validation: | | 0/? [00:00<?, ?it/s]

Metric val_loss improved by 0.074 >= min_delta = 0.0. New best score: 1.705

Validation: | | 0/? [00:00<?, ?it/s]

Metric val_loss improved by 0.051 >= min_delta = 0.0. New best score: 1.653

Validation: | | 0/? [00:00<?, ?it/s]

Metric val_loss improved by 0.141 >= min_delta = 0.0. New best score: 1.512

Validation: | | 0/? [00:00<?, ?it/s]

Metric val_loss improved by 0.030 >= min_delta = 0.0. New best score: 1.482

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Monitored metric val_loss did not improve in the last 3 records. Best score: 1.482. Signaling Trainer to stop.

```

In [15]: # Actual Model (no Trainer) to Predict!
pred_output = model.predict(test_input_ids_tensors, test_attention_mask_tensors)

# Show Tensor Output Probailities!
print(pred_output)

```

```

tensor([[0.6053, 0.0123, 0.2501, ..., 0.0112, 0.0404, 0.0141],
        [0.1246, 0.0130, 0.0440, ..., 0.0193, 0.0788, 0.0202],
        [0.0800, 0.0165, 0.0297, ..., 0.0445, 0.0749, 0.0389],
        ...,
        [0.0145, 0.0898, 0.0216, ..., 0.3688, 0.0394, 0.4252],
        [0.0450, 0.0356, 0.0283, ..., 0.2031, 0.0684, 0.4416],
        [0.0132, 0.0727, 0.0182, ..., 0.3354, 0.0376, 0.4805]])

```

```

In [16]: # Get Predictions
y_preds = np.argmax(pred_output, axis=1)

```



```
# Get the F1-Score and Accuracy of the Correlation Matrix Details
f1 = f1_score(test_labels_tensors, y_preds, average="weighted")
acc = accuracy_score(test_labels_tensors, y_preds)

print(f"Acc: {acc}")
print(f"F1-Score: {f1}")
print("|")
print(f'Output Prediction:{y_preds.shape}')
print(f'Pred\n: {y_preds}')
print(f'Label\n: {test_labels_tensors}')
```

Acc: 0.38095238095238093

F1-Score: 0.3472112521017685

|

Output Prediction:torch.Size([147])

Pred

```
: tensor([0, 3, 3, 3, 3, 3, 0, 0, 6, 0, 2, 3, 0, 3, 3, 6, 6, 3, 3, 3, 6, 6, 0, 2,
          3, 3, 6, 6, 6, 0, 0, 0, 3, 6, 6, 3, 0, 3, 2, 6, 2, 6, 3, 3, 2, 0, 6, 3,
          6, 3, 3, 6, 3, 4, 6, 3, 3, 6, 0, 0, 3, 3, 3, 2, 0, 6, 0, 6, 6, 2, 0, 2,
          0, 2, 0, 6, 6, 6, 0, 6, 6, 6, 2, 6, 6, 4, 3, 6, 2, 3, 6, 3, 6, 6, 4, 0,
          6, 0, 3, 3, 6, 0, 3, 4, 6, 4, 6, 3, 4, 6, 6, 6, 6, 3, 0, 6, 3, 6, 3, 3,
          3, 6, 0, 6, 3, 3, 3, 6, 6, 6, 0, 6, 3, 0, 0, 6, 6, 6, 4, 3, 0, 0, 3, 3,
          6, 6, 6])
```

Label

```
: tensor([3, 3, 6, 3, 3, 6, 0, 2, 4, 0, 2, 3, 0, 5, 0, 4, 1, 3, 4, 3, 3, 5, 2, 0,
          2, 3, 2, 4, 4, 3, 0, 0, 2, 0, 1, 3, 0, 6, 0, 4, 5, 4, 3, 3, 0, 5, 0, 4,
          3, 3, 3, 2, 6, 0, 6, 2, 3, 4, 2, 3, 3, 0, 3, 2, 2, 5, 6, 3, 6, 2, 3, 0,
          6, 2, 0, 1, 4, 0, 6, 6, 6, 6, 4, 3, 6, 4, 0, 6, 2, 2, 4, 6, 4, 6, 6, 5,
          5, 6, 4, 3, 6, 6, 3, 2, 3, 4, 4, 3, 6, 1, 3, 3, 4, 0, 3, 6, 3, 4, 3, 0,
          6, 6, 5, 4, 4, 3, 6, 6, 4, 2, 0, 6, 6, 6, 6, 1, 3, 6, 4, 3, 4, 0, 3, 6,
          6, 5, 6])
```

```
In [17]: # Create Confusion Matrix
conf_matrix = confusion_matrix(test_labels_tensors, y_preds)

# Show Details
print(conf_matrix)
print()

# Get Class Names based on teh Label Idx
sorted_df = df.sort_values(by='label').drop_duplicates(subset=['emotion_en'])
class_names = sorted_df['emotion_en'].tolist()
# Something like this: class_names = ['Happiness', 'Hate', 'Love', 'Neutral', 'Sadness', 'Surprise', 'Worry']

plt.figure(figsize=(20, 8))

# Create Seaborn Heatmap for Display Confusion Matrix
plt.subplot(121)
sns.heatmap(conf_matrix,
            annot=True,
            fmt='g',
            cmap='Blues',
            linewidth=1,
            xticklabels=class_names,
            yticklabels=class_names)

# Setup Title and Axis Lables
plt.title('Emotion Confusion Matrix\n(Test Dataset)\n',
          fontweight='bold',
          fontsize='18',
          horizontalalignment='center')

plt.xlabel('\nPredicted Labels',
          fontweight='bold',
          fontsize='18',
          horizontalalignment='center')

plt.ylabel('True Labels\n',
          fontweight='bold',
          fontsize='18')

plt.subplot(122)
sns.histplot(data=df, x='emotion_en')

# Setup Title and Axis Lables
plt.title('Emotion Histogram\n',
          fontweight='bold',
          fontsize='18',
          horizontalalignment='center')

plt.xlabel('\nEmotions',
          fontweight='bold',
```

```

        fontsize='18',
        horizontalalignment='center')

plt.ylabel('Counts / Frequency',
           fontweight='bold',
           fontsize='18')

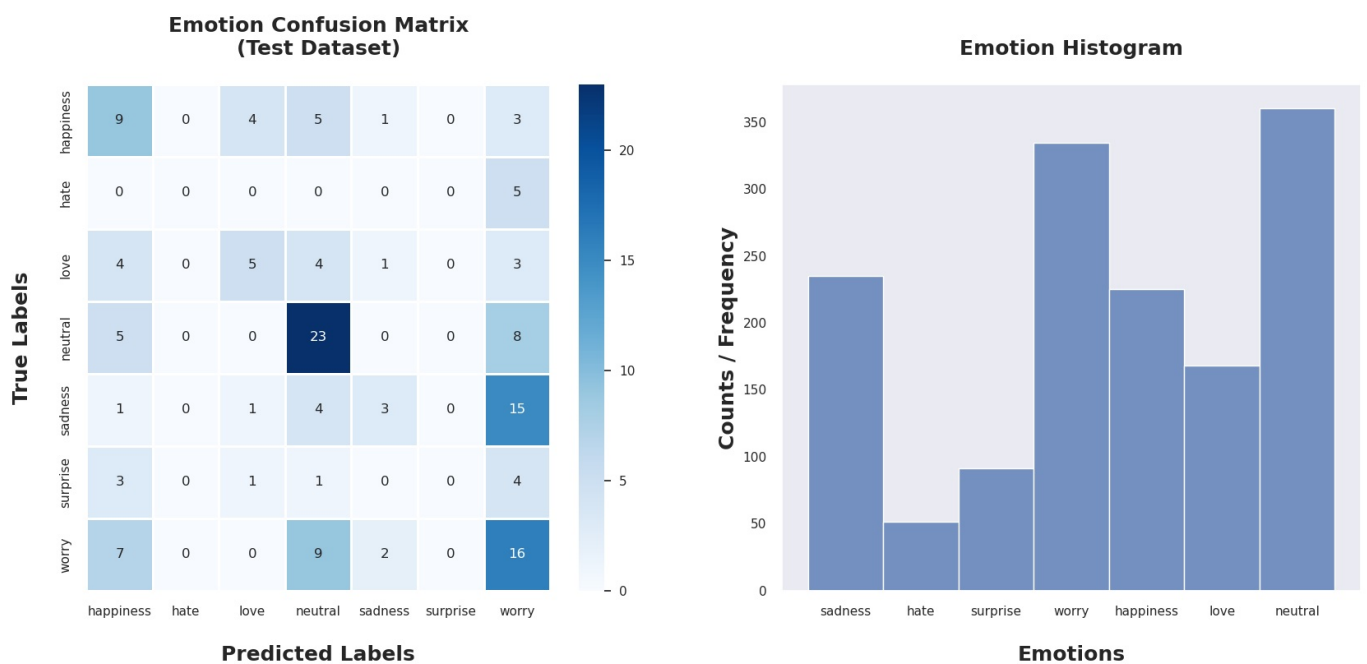
# Save & Show the Plot
plt.savefig('fig_confusion_matrix.png')
plt.show()

```

```

[[ 9  0  4  5  1  0  3]
 [ 0  0  0  0  0  0  5]
 [ 4  0  5  4  1  0  3]
 [ 5  0  0 23  0  0  8]
 [ 1  0  1  4  3  0 15]
 [ 3  0  1  1  0  0  4]
 [ 7  0  0  9  2  0 16]]

```



Display Validation & Testing Results

```

In [18]: trainer.validate(model, dataloaders=val_dataloader)
         trainer.test(model, test_dataloader)

```

```

LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
SLURM auto-queueing enabled. Setting signal handlers.
Validation: |          | 0/? [00:00<?, ?it/s]

```

Validate metric	DataLoader 0
val_acc	0.39726027846336365
val_loss	1.5434670448303223

```

LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
SLURM auto-queueing enabled. Setting signal handlers.
Testing: |          | 0/? [00:00<?, ?it/s]

```

Test metric	DataLoader 0
test_acc	0.380952388048172
test_loss	1.7607795000076294

```

Out[18]: [{'test_loss': 1.7607795000076294, 'test_acc': 0.380952388048172}]

```

Plot Training Parameters

```

In [19]: results = pd.read_csv(csv_logger.log_dir + "/metrics.csv")
         results

```

Out[19]:

	val_loss	step	epoch	val_acc	train_loss	lr	train_acc	test_loss	test_acc
0	1.778817	117	0	0.294521	NaN	NaN	NaN	NaN	NaN
1	NaN	117	0	NaN	1.853640	0.000005	0.236550	NaN	NaN
2	1.704656	235	1	0.328767	NaN	NaN	NaN	NaN	NaN
3	NaN	235	1	NaN	1.742712	0.000005	0.316823	NaN	NaN
4	1.653244	353	2	0.417808	NaN	NaN	NaN	NaN	NaN
5	NaN	353	2	NaN	1.625816	0.000005	0.381725	NaN	NaN
6	1.511819	471	3	0.397260	NaN	NaN	NaN	NaN	NaN
7	NaN	471	3	NaN	1.509168	0.000005	0.444065	NaN	NaN
8	1.481664	589	4	0.390411	NaN	NaN	NaN	NaN	NaN
9	NaN	589	4	NaN	1.391266	0.000005	0.491887	NaN	NaN
10	1.509674	707	5	0.383562	NaN	NaN	NaN	NaN	NaN
11	NaN	707	5	NaN	1.279503	0.000005	0.536294	NaN	NaN
12	1.506783	825	6	0.376712	NaN	NaN	NaN	NaN	NaN
13	NaN	825	6	NaN	1.177210	0.000004	0.579846	NaN	NaN
14	1.543467	943	7	0.397260	NaN	NaN	NaN	NaN	NaN
15	NaN	943	7	NaN	1.083147	0.000004	0.591802	NaN	NaN
16	1.543467	944	8	0.397260	NaN	NaN	NaN	NaN	NaN
17	NaN	944	8	NaN	NaN	NaN	NaN	1.76078	0.380952

In [20]:

```
# Set Chart Sizing
plt.figure(figsize=(15, 12))

# Best EPOCH
BEST_EPOCH = 4

# Loss Plot
plt.subplot(221)
plt.plot(results["epoch"][np.logical_not(np.isnan(results["train_loss"]))],
         results["train_loss"][np.logical_not(np.isnan(results["train_loss"]))],
         label="Training")

plt.plot(results["epoch"][np.logical_not(np.isnan(results["val_loss"]))],
         results["val_loss"][np.logical_not(np.isnan(results["val_loss"]))],
         label="Validation")

plt.legend()
plt.title('Training\nValidation Loss',
         fontweight='bold',
         fontsize='18',
         horizontalalignment='center')

plt.xlabel('Epochs',
         fontweight='bold',
         fontsize='18',
         horizontalalignment='center')

plt.ylabel('Loss',
         fontweight='bold',
         fontsize='18')

# Show Epoch that was best!
plt.axvline(x=BEST_EPOCH, color='r', linestyle='--')

# Accuracy Plot
plt.subplot(222)
plt.plot(results["epoch"][np.logical_not(np.isnan(results["train_acc"]))],
         results["train_acc"][np.logical_not(np.isnan(results["train_acc"]))],
         label="Training")

plt.plot(results["epoch"][np.logical_not(np.isnan(results["val_acc"]))],
         results["val_acc"][np.logical_not(np.isnan(results["val_acc"]))],
         label="Validation")

plt.legend()
plt.title('Training\nValidation Accuracy',
         fontweight='bold',
         fontsize='18',
```

```
        horizontalalignment='center')

plt.xlabel('Epochs',
           fontweight='bold',
           fontsize='18',
           horizontalalignment='center')

plt.ylabel('Accuracy',
           fontweight='bold',
           fontsize='18')

# Show Epoch that was best!
plt.axvline(x=BEST_EPOCH, color='r', linestyle='--')

# LR Plot
plt.subplot(223)
plt.plot(results["epoch"][np.logical_not(np.isnan(results["lr"]))],
         results["lr"][np.logical_not(np.isnan(results["lr"]))],
         label="Learning Rate")

plt.legend()
plt.title('Learning Rate Decay',
         fontweight='bold',
         fontsize='18',
         horizontalalignment='center')




plt.xlabel('Epochs',
         fontweight='bold',
         fontsize='18',
         horizontalalignment='center')

plt.ylabel('Learning Rate',
         fontweight='bold',
         fontsize='18')

# Show Epoch that was best!
plt.axvline(x=BEST_EPOCH, color='r', linestyle='--')

plt.tight_layout(h_pad=1.5)

# Save & Show the Plot
plt.savefig('fig_training.png')
plt.show()
```



In []: